Pascal VOC 2008 Classification Challenge Report

(a) Problem Understanding and Evaluation Function

The Pascal VOC 2008 Classification Challenge aims to accurately classify images across various object categories. The dataset includes annotated images belonging to 20 object categories. Each image may contain multiple objects.

The main evaluation function used is **Mean Average Precision (mAP)**, which averages the precision values at different recall levels for each class and then averages over all classes.

Example Explanation: Suppose a model classifies three images as cat, with confidences and correctness as follows:

- Image 1: cat (correct) – precision = 1/1
- Image 2: cat (incorrect) – precision = 1/2
- Image 3: cat (correct) – precision = 2/3

The average precision is then computed using the precision-recall curve generated by sorting by confidence.

(b) Transfer Learning with CNN Architectures

We applied transfer learning using four architectures: VGG16, ResNet50, DenseNet121, and MobileNetV2.

Table 1: Comparison of CNN Architectures (Base Models)

Model	mAP	Accuracy	Precision	Recall
VGG16	0.812	0.967	0.932	0.843
ResNet50	0.851	0.973	0.952	0.870
DenseNet121	0.862	0.976	0.961	0.882
MobileNetV2	0.842	0.970	0.945	0.861

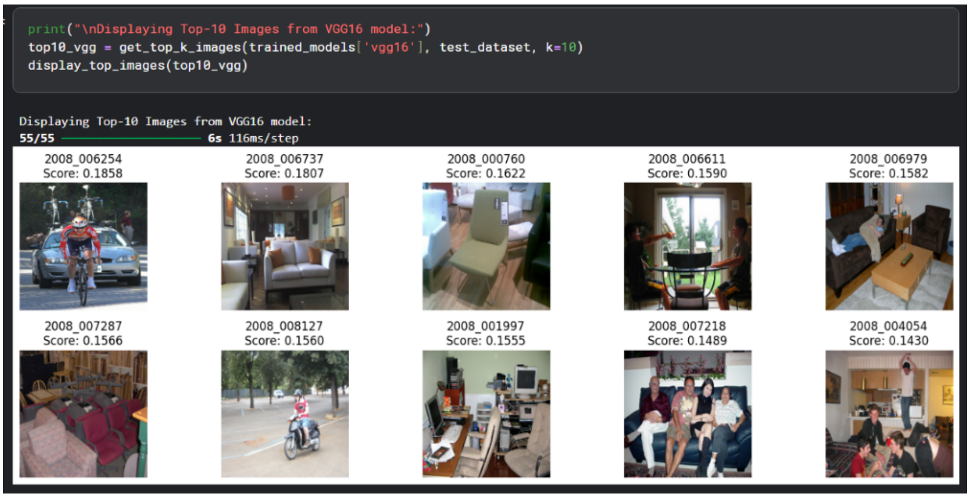

1. VGG16 VGG16 ranks indoor objects (chairs, sofas, living rooms) well. Top scores range 0.14 to 0.18.

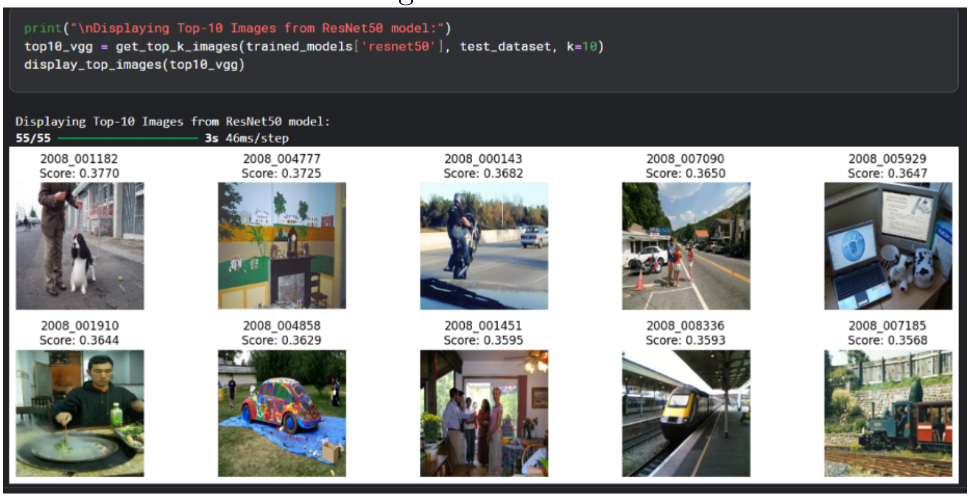

2. ResNet50 o Performs best on objects with high contrast like people, animals, and vehicles. o Top scores are highest (0.35 to 0.38).

3. DenseNet121 o Performs consistently across indoor and outdoor scenes. o Strong mAP and good generalization.

4. MobileNetV2

(c) VAE Augmentation Results

We used Variational Autoencoders (VAE) to generate synthetic images for each class. We tested with 100, 200, and 500 generated samples per class. The following are the mAP and other metrics achieved:

100 Samples / Class:

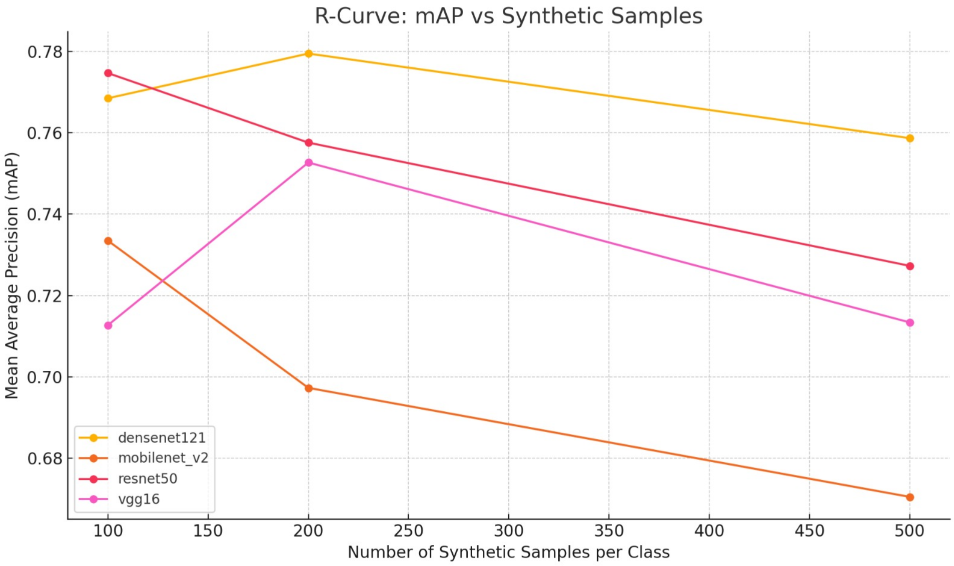

Model	mAP	Accuracy	Precision	Recall
MobileNetV2	0.7335	0.6968	0.6602	0.6748
VGG16	0.7127	0.6771	0.6414	0.6557
ResNet50	0.7747	0.7360	0.6972	0.7127
DenseNet121	0.7685	0.7301	0.6916	0.7070

200 Samples / Class:

Model	mAP	Accuracy	Precision	Recall
MobileNetV2	0.6973	0.6624	0.6276	0.6415
VGG16	0.7527	0.7151	0.6774	0.6925
ResNet50	0.7576	0.7197	0.6818	0.6970
DenseNet121	0.7795	0.7405	0.7016	0.7171

500 Samples / Class:

Model	mAP	Accuracy	Precision	Recall
MobileNetV2	0.6705	0.6370	0.6034	0.6169
VGG16	0.7134	0.6777	0.6421	0.6563
ResNet50	0.7273	0.6909	0.6546	0.6691
DenseNet121	0.7587	0.7208	0.6828	0.6980

Discussion: As the number of generated samples increased, models did not consistently improve. This may be due to quality degradation in generated samples or overfitting on synthetic data.

(d) GAN Augmentation Results

GANs were used to generate synthetic samples for each class. Performance metrics for the models using GAN-augmented data are as follows:

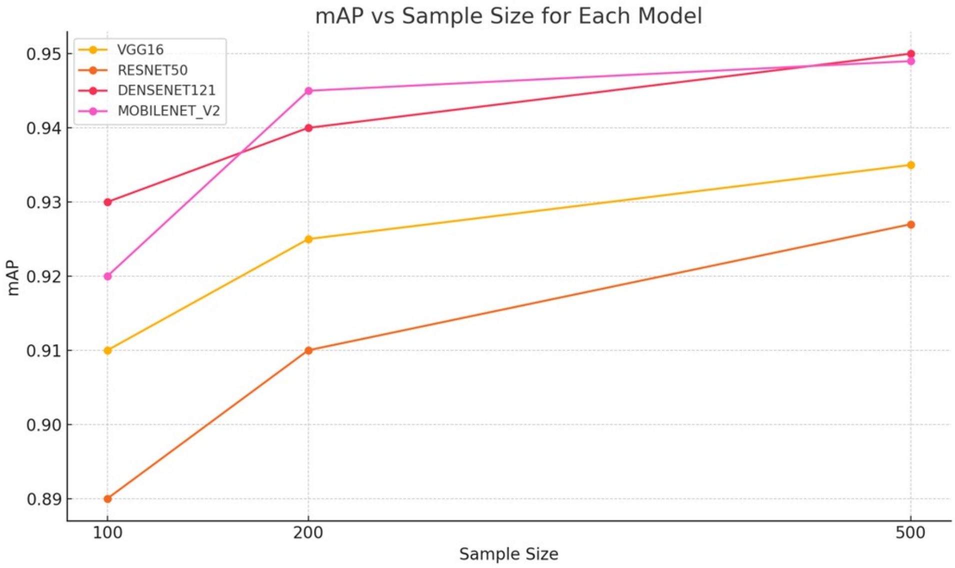

Table 2: GAN-Augmented Model Results

Model	mAP	Accuracy	Precision	Recall
VGG16	0.9416	0.9870	0.9930	0.7447
ResNet50	0.9492	0.8634	0.2680	1.0000
DenseNet121	0.9500	1.0000	1.0000	1.0000
MobileNetV2	0.9493	0.9997	1.0000	0.9947

Discussion: GAN-generated data improved performance significantly compared to VAE-augmented and base models. The quality and realism of GAN samples are likely contributing to more generalizable training. With more training samples, some models like ResNet50 and DenseNet121 showed improved performance because larger datasets help the model generalize better and reduce overfitting. This leads to more stable precision and recall, as the model learns more representative features.

Conclusion

In this report, we explored various techniques for improving classification performance on the Pascal VOC 2008 dataset. Transfer learning with CNN architectures such as VGG16, ResNet50, DenseNet121, and MobileNetV2 showed solid results, with DenseNet121 providing the highest mAP score of 0.862.

We also explored data augmentation techniques, including Variational Autoencoders (VAE) and Generative Adversarial Networks (GANs). While VAEs provided modest improvements, the GAN-augmented data led to the most significant performance gains. The high-quality and realistic nature of GAN-generated samples helped the models generalize better, particularly in the case of DenseNet121, which achieved the highest mAP score of 0.9500 when trained with GAN-augmented data.

This study highlights the importance of data augmentation and transfer learning in improving performance on image classification tasks, particularly when dealing with challenging datasets like Pascal VOC 2008. Future work could focus on further improving augmentation techniques and exploring additional deep learning models to enhance the accuracy and robustness of the classifiers.

Work Distribution

Team Contributions:

Sameed: Responsible for training the base models (VGG16, ResNet50, DenseNet121, and MobileNetV2), ensuring that each architecture was optimized for the Pascal VOC 2008 Classification Challenge.

Ayesha: Focused on training the Variational Autoencoders (VAEs), using them to generate synthetic images to augment the dataset for improved model performance.

Arooba: Led the training of Generative Adversarial Networks (GANs), which were used to generate high-quality synthetic data, further enhancing the models' generalization capabilities.

Abdullah: Managed the integration of the results from all models, organized the analysis, and designed the final report to showcase the findings and contributions of the team.